



ITDS 122 Data Structures and Algorithm Analysis

Lecture 14: Graph Algorithms

Asst. Prof. Dr. Ananta Srisuphab
2/2024



Teaching Lecture Hour : 30 Lab Hour : 30

- Recalling Python Coding
- Object-Oriented Programming with Python
- Algorithm Analysis
- Recursion
- Array-Based Sequences
- Stacks, Queues, Deques
- Linked Lists
- Review & checkpoint [Midterm]
- Trees and Search Trees *
- Maps and Hash Tables
- Graph Algorithms *
- Sorting
- Review & checkpoint [Final]

Graph Algorithms

- Transitive Closure
- Directed Acyclic Graphs (DAGs)
- Shortest Paths
- Minimum Spanning Tree (MST)

Transitive Closure

Introduction to Transitive Closure

- Goal: Determine reachability in a directed graph
- If vertex u can reach vertex v , then v is in the transitive closure of u
- Applications:
 - Routing & network analysis
 - Scheduling systems
 - Analyzing dependencies

Definition of Transitive Closure

- Transitive closure G^* of a directed graph G :
 - Same vertex set as G
 - Edge (u, v) exists in G^* iff there is a path from u to v in G
- Can include self-loops (path from u to u)

Motivation for Precomputing

- DFS/BFS answers reachability for a single source
- Inefficient for multiple queries
- Precompute transitive closure for constant-time reachability checks
- Efficient for dense graphs or matrix-based representations

Floyd-Warshall Algorithm

- Iteratively add paths by checking intermediate vertices
- Let vertices be v_1, v_2, \dots, v_n
- For each k , add (i, j) if $i \rightarrow k$ and $k \rightarrow j$ exist
- Key idea: path $i \rightarrow k \rightarrow j$ implies $i \rightarrow j$

Pseudocode

Algorithm FloydWarshall(\vec{G}):

Input: A directed graph \vec{G} with n vertices

Output: The transitive closure \vec{G}^* of \vec{G}

let v_1, v_2, \dots, v_n be an arbitrary numbering of the vertices of \vec{G}

$\vec{G}_0 = \vec{G}$

for $k = 1$ to n **do**

$\vec{G}_k = \vec{G}_{k-1}$

for all i, j in $\{1, \dots, n\}$ with $i \neq j$ and $i, j \neq k$ **do**


if both edges (v_i, v_k) and (v_k, v_j) are in \vec{G}_{k-1} **then**

add edge (v_i, v_j) to \vec{G}_k (if it is not already present)


return \vec{G}_n

Python Implementation

```
4 class Graph:
5     def __init__(self):
6         self.adj = {}
7
8     def add_vertex(self, v):
9         if v not in self.adj:
10            self.adj[v] = set()
11
12    def add_edge(self, u, v):
13        self.add_vertex(u)
14        self.add_vertex(v)
15        self.adj[u].add(v)
16
17    def vertices(self):
18        return list(self.adj.keys())
19
20    def get_edge(self, u, v):
21        return v if v in self.adj.get(u, set()) else None
22
23    def insert_edge(self, u, v):
24        self.adj[u].add(v)
25
26    def __str__(self):
27        return "\n".join(f"{u} -> {sorted(vs)}" for u, vs in self.adj.items())
28
```



(a)



(f)

```

29 def floyd_warshall(g):
30     """Return a new graph that is the transitive closure of g."""
31     closure = deepcopy(g)
32     verts = list(closure.vertices())
33     n = len(verts)
34
35     for k in range(n):
36         for i in range(n):
37             if i != k and closure.get_edge(verts[i], verts[k]) is not None:
38                 for j in range(n):
39                     if i != j != k and closure.get_edge(verts[k], verts[j]) is not None:
40                         if closure.get_edge(verts[i], verts[j]) is None:
41                             closure.insert_edge(verts[i], verts[j])
42
43     return closure
44 
```

```

45 #
46 g = Graph()
47 g.add_edge("LAX", "ORD")
48 g.add_edge("DFW", "LAX")
49 g.add_edge("DFW", "SFO")
50 g.add_edge("DFW", "ORD")
51 g.add_edge("ORD", "DFW")
52 g.add_edge("MIA", "LAX")
53 g.add_edge("MIA", "DFW")
54 g.add_edge("JFK", "DFW")
55 g.add_edge("JFK", "MIA")
56 g.add_edge("JFK", "BOS")
57 g.add_edge("BOS", "JFK")
58
59
60
61 if __name__ == "__main__":
62     print("Original Graph:")
63     print(g)
64
65     closure = floyd_warshall(g)
66     print("\nTransitive Closure:")
67     print(closure)
68 
```



```



Original Graph:
LAX -> ['ORD']
ORD -> ['DFW']
DFW -> ['LAX', 'ORD', 'SFO']
SFO -> []
MIA -> ['DFW', 'LAX']
JFK -> ['BOS', 'DFW', 'MIA']
BOS -> ['JFK']


Transitive Closure:
LAX -> ['DFW', 'ORD', 'SFO']
ORD -> ['DFW', 'LAX', 'SFO']
DFW -> ['LAX', 'ORD', 'SFO']
SFO -> []
MIA -> ['DFW', 'LAX', 'ORD', 'SFO']
JFK -> ['BOS', 'DFW', 'LAX', 'MIA', 'ORD', 'SFO']
BOS -> ['DFW', 'JFK', 'LAX', 'MIA', 'ORD', 'SFO']

```


Visual Walkthrough







(e)



(f)

Summary

- Transitive closure = precomputed reachability
- Floyd-Warshall is conceptually simple and easy to implement
- Use for dense graphs or matrix-based representations
- Alternatives: DFS/BFS for sparse graphs (lower time but repeated calls)

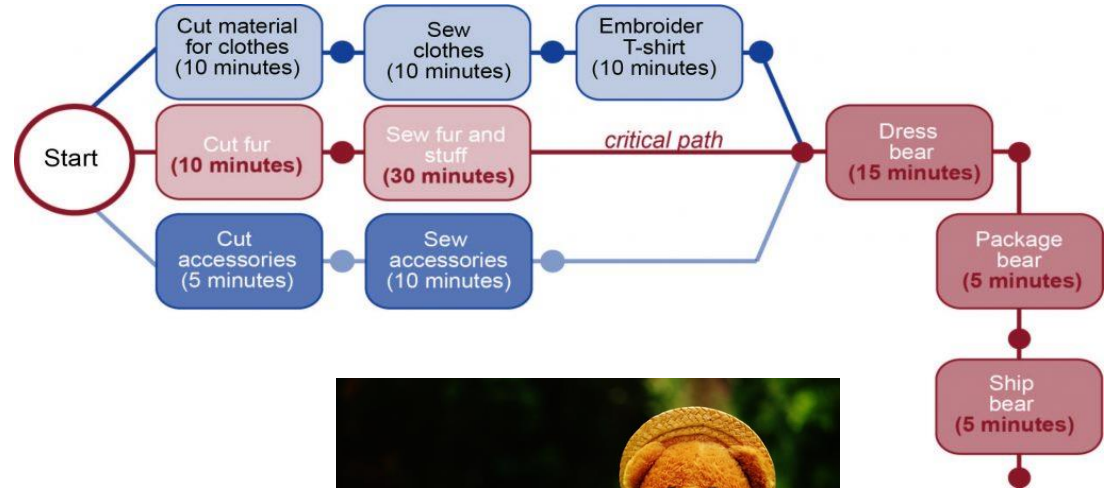
Directed Acyclic Graphs (DAGs)

Introduction to Directed Acyclic Graphs (DAGs)

- A Directed Acyclic Graph (DAG) is a directed graph with no cycles.
- Common applications:
 - Course prerequisites
 - Task scheduling
 - Inheritance hierarchies in OOP
- Important property: DAGs can be topologically ordered.

Motivation Example

- To manage a large project, break it into smaller tasks.
- Some tasks depend on others (e.g., foundation before walls).
- Represent tasks as vertices, dependencies as directed edges.
- If feasible, the graph must be a DAG.





Topological Ordering

- A topological ordering of a DAG:
 - A sequence of vertices such that for every edge (u, v) , u comes before v .
- Proposition: A graph has a topological order iff it is acyclic.

Visual Example

- Two possible topological orders for the same DAG:
 - Different valid sequences based on constraint satisfaction.
- Key point: multiple valid topological sorts may exist.



Topological Sorting Algorithm : Kahn's Algorithm

1. Initialize:

- Count incoming edges (in-degree)
- Add nodes with 0 in-degree to a ready list

2. While ready is not empty:

- Remove node u from ready
- Append u to topo
- Decrease in-degree of u 's neighbors
- Add new nodes with 0 in-degree to ready


Python Implementation

```
3 def topological_sort(g):
4     """
5     Return a list of vertices of directed acyclic graph g in topological order.
6     If graph g has a cycle, the result will be incomplete.
7     """
8     topo = []                # list of vertices placed in topological order
9     ready = []              # list of vertices that have no remaining constraints
10    incount = {}             # keep track of in-degree for each vertex
11
12    for u in g:
13        incount[u] = 0
14
15    for u in g:
16        for v in g[u]:
17            incount[v] += 1
18
19    for u in g:
20        if incount[u] == 0:
21            ready.append(u)
22
23    while ready:
24        u = ready.pop()
25        topo.append(u)
26        for v in g[u]:
27            incount[v] -= 1
28            if incount[v] == 0:
29                ready.append(v)
30
31    return topo
```


```
34 graph = {
35     'A': ['C', 'D'],
36     'B': ['D', 'F'],
37     'C': ['D', 'E', 'H'],
38     'D': ['E', 'F'],
39     'E': ['G'],
40     'F': ['G', 'H'],
41     'G': ['H'],
42     'H': []
43 }
44
45 if __name__ == "__main__":
46     ordering = topological_sort(graph)
47     print("Topological Ordering:", ordering)
48
```

Example Walkthrough

- Highlight the step-by-step updates:
 - Track incout changes
 - Highlight which node is selected next
 - Show dashed edges as processed
- Visualize how nodes are added to final order.



(a)




(b)

Selected Node: A (first in the ready list)
Topological Order: [A]
Ready List: [B]

Update in-degrees for neighbors of A:

- A → C → $\text{incount}[C] = 1 \rightarrow 1 \rightarrow 0$
- A → D → $\text{incount}[D] = 1 \rightarrow 3 \rightarrow 2$

Add C to ready (since $\text{incount}[C] == 0$)



(c)

Selected Node: C
Ready List: [B]

Update in-degrees for C's neighbors:


- C → D → $\text{incount}[D] = 1 \rightarrow 2 \rightarrow 1$
- C → E → $\text{incount}[E] = 1 \rightarrow 1 \rightarrow 0$
- C → H → $\text{incount}[H] = 1 \rightarrow 3 \rightarrow 2$

Add E to ready (because $\text{incount}[E] == 0$)


Ready List: [A, B]
Topological Order: []

Ready List: [B, C]
Topological Order: [A]


Ready List: [B, E]
Topological Order: [A, C]




(d)




(e)




(f)



(g)



(h)



(i)

Performance Analysis

- Let n = number of vertices, m = number of edges
- Time complexity: $O(n + m)$
- Space complexity: $O(n)$
- Efficient for large sparse graphs

Detecting Cycles

- If topo does not include all nodes:
 - The graph contains a cycle
 - Some nodes remain with non-zero in-degree

Summary

- DAGs are fundamental in scheduling and dependency resolution
- Topological sorting is a key algorithm on DAGs
- Python implementation uses Kahn's algorithm efficiently
- Always verify acyclicity to ensure a valid ordering

Shortest Paths

Introduction to Shortest Paths

- Goal: Find the shortest (minimum-weight) path between vertices
- Applications:
 - Navigation systems
 - Routing in computer networks
 - Task scheduling
- Breadth-First Search (BFS) only works when all edge weights are equal

Weighted Graphs

- A weighted graph assigns a numeric weight $w(e)$ to each edge e
- For edge (u, v) , use notation: $w(u, v) = w(e)$
- Weights can represent:
 - Distance
 - Cost
 - Time

Path and Distance Definitions

- A path from u to v is a sequence of edges
- The length of a path = sum of its edge weights
- The shortest path from u to v is a path with minimum total weight
- If no path exists: $d(u, v) = \infty$

Valid Weights for Shortest Path Algorithms

- Must use nonnegative edge weights for Dijkstra
- Negative-weight edges allow cycles that reduce cost arbitrarily
- Algorithms must avoid revisiting same nodes with lower cost in cycles

Dijkstra's Algorithm Overview

- Greedy algorithm for single-source shortest paths
- Always picks vertex with smallest tentative distance ($D[v]$)
- Expands a “cloud” of known shortest paths from the source
- Relaxes edges from each newly added vertex

Edge Relaxation Explained

- For each edge (u, v) :

```
if  $D[v] > D[u] + w(u, v)$ :  
     $D[v] = D[u] + w(u, v)$ 
```

- Ensures shortest known distance is improved
- $D[v]$ values always shrink or stay the same

Dijkstra Pseudocode

Algorithm ShortestPath(G, s):

Input: A weighted graph G with nonnegative edge weights, and a distinguished vertex s of G .

Output: The length of a shortest path from s to v for each vertex v of G .

Initialize $D[s] = 0$ and $D[v] = \infty$ for each vertex $v \neq s$.

Let a priority queue Q contain all the vertices of G using the D labels as keys.

while Q is not empty **do**

 {pull a new vertex u into the cloud}

$u =$ value returned by $Q.\text{remove_min}()$

for each vertex v adjacent to u such that v is in Q **do**

 {perform the *relaxation* procedure on edge (u, v) }



if $D[u] + w(u, v) < D[v]$ **then**



$D[v] = D[u] + w(u, v)$


 Change to $D[v]$ the key of vertex v in Q .

return the label $D[v]$ of each vertex v


Example Execution









(e)




(f)




(g)



(h)



(i)



(j)

Reconstructing Shortest Paths

- Dijkstra gives shortest distance, but not actual path
- To reconstruct:
 - For each $v \neq s$, find edge (u, v) such that $D[u] + w(u, v) = D[v]$
- Builds the shortest-path tree

Summary

- Dijkstra's algorithm solves the single-source shortest path problem
- Assumes nonnegative weights
- Builds both distance map and shortest-path tree
- Efficient and correct with proper data structures

Minimum Spanning Tree (MST)

Minimum Spanning Tree (MST)

- Connect all vertices of a weighted, undirected graph
- Use minimum total edge weight
- Real-world: Lay out cable to connect offices with minimal cost

Problem Definition

- Given undirected graph $G = (V, E)$ with weights $w(e)$
- Find tree $T \subseteq E$ that spans all V with minimal:

$$w(T) = \sum w(u, v) \quad \text{for all } (u, v) \text{ in } T$$

- T is a Minimum Spanning Tree (MST)

Prim-Jarník Algorithm - Concept

- Grow MST from an arbitrary start vertex
- Use a greedy strategy: always add lightest edge to expand MST
- Very similar structure to Dijkstra's Algorithm

Prim-Jarník Pseudocode

Algorithm PrimJarnik(G):

Input: An undirected, weighted, connected graph G with n vertices and m edges

Output: A minimum spanning tree T for G

Pick any vertex s of G

$D[s] = 0$

for each vertex $v \neq s$ **do**

$D[v] = \infty$

Initialize $T = \emptyset$.

Initialize a priority queue Q with an entry $(D[v], (v, \text{None}))$ for each vertex v , where $D[v]$ is the key in the priority queue, and (v, None) is the associated value.

while Q is not empty **do**

$(u, e) = \text{value returned by } Q.\text{remove_min}()$

 Connect vertex u to T using edge e .

for each edge $e' = (u, v)$ such that v is in Q **do**

 {check if edge (u, v) better connects v to T }

if $w(u, v) < D[v]$ **then**


$D[v] = w(u, v)$

 Change the key of vertex v in Q to $D[v]$.


 Change the value of vertex v in Q to (v, e') .

return the tree T


MST Illustration




(a)




(b)




(c)





(d)




(e)




(f)





(i)



(j)

Kruskal's Algorithm for MST

- Greedy algorithm for Minimum Spanning Tree (MST)
- Builds MST by adding lightest edges one-by-one
- Keeps growing a forest (set of trees), until connected

Kruskal's Algorithm Concept

- Start: All vertices are individual clusters (singleton sets)
- Sort edges by weight (min \rightarrow max)
- For each edge (u, v) :
 - If u and v are in different clusters \rightarrow add edge to MST
 - Merge their clusters
 - If u and v are already connected \rightarrow skip (to avoid cycle)

Kruskal's Pseudocode

Algorithm Kruskal(G):

Input: A simple connected weighted graph G with n vertices and m edges

Output: A minimum spanning tree T for G

for each vertex v in G do

Define an elementary cluster $C(v) = \{v\}$.

Initialize a priority queue Q to contain all edges in G , using the weights as keys.

$T = \emptyset$	$\{T \text{ will ultimately contain the edges of the MST}\}$
-----------------	--

while T has fewer than $n - 1$ edges do

(u, v) = value returned by `Q.remove_min()`

Let $C(u)$ be the cluster containing u , and let $C(v)$ be the cluster containing v .



if $C(u) \neq C(v)$ then



Add edge (u, v) to T .


Merge $C(u)$ and $C(v)$ into one cluster.

return tree T


Kruskal's MST algorithm Illustration












(e)




(f)







(k)



(l)

